



Mixed decode requests

one long request and several short requests



Request-level scheduling

whole request is bound to one worker



Token/request-level scheduling

split a long request while short requests stay local

